

Review of Rashid & Nanzer, “Frequency and Phase Synchronization in Distributed Antenna Arrays Based on Consensus Averaging and Kalman Filtering”

IEEE Transactions on Wireless Communications, vol. 22, no. 4, pp. 2789–2802, April 2023.
DOI 10.1109/TWC.2022.3213788

Chandler Stevenson

July 28, 2026

1 Summary

The paper addresses joint frequency and phase alignment of N oscillators in an open-loop coherent distributed array—one with no feedback from the beamforming destination. Each node locally broadcasts its oscillator frequency and phase to its neighbors over an undirected connectivity graph. In the baseline algorithm, *decentralized frequency and phase consensus* (DFPC), every node iteratively replaces its state with a weighted (Metropolis–Hastings) average of its neighbors’; the doubly stochastic mixing matrix $\mathbf{W}$ drives all nodes to the arithmetic mean of their initial states, with convergence speed and residual error governed by the second-largest eigenvalue modulus λ_2 .

Errors are injected at each iteration from a practical statistical model: oscillator frequency drift from a two-term Allan-deviation law $\sigma_f = f_c \sqrt{\beta_1/T + \beta_2 T}$ (white FM plus random-walk FM, $\beta_1 = \beta_2 = 5 \times 10^{-19}$ for a quartz TCXO), intra-interval drift, phase jitter drawn per interval from the oscillator’s integrated phase-noise power ($A = -53.46$ dBc $\Rightarrow \sigma_\theta = 2.7$ mrad), and frequency/phase estimation errors at their Cramér–Rao bounds for an L -sample observation at a given SNR. The central analytical result (their Eq. 27) is a closed-form steady-state total phase error for DFPC on sparsely connected graphs: a root-sum-of-squares of five terms — oscillator drift over the update interval ($2\pi\sigma_f T$), intra-interval drift, phase jitter, and the two estimation errors.

To reduce the residual at short update intervals, the paper integrates a per-node two-state Kalman filter (state $[f_n, \theta_n]$, process noise from the drift/jitter statistics, measurement noise at the CRLBs) whose posterior means are then consensus-averaged (KF-DFPC). Monte Carlo results ($f_c = 1$ GHz, $f_s = 10$ MHz, $T = 0.1$ ms typical, 10^3 trials) show: (i) KF-DFPC converges in dramatically fewer iterations than DFPC on sparse graphs (e.g., 2 vs. 16 iterations at $N=100$, connectivity 0.05; 4 vs. 158 at $N=65$); (ii) it outperforms diffusion LMS, diffusion KF, and the Kalman consensus information filter in various regimes; (iii) its converged total phase error is essentially independent of SNR, unlike DFPC, whose error at low SNR is estimation-dominated; and (iv) sweeping the update interval reveals an optimum ($T \approx 20$ ms at their parameters) where drift-driven and estimation-driven errors balance—DFPC degrades on either side, while KF-DFPC holds the theoretical floor for $T \lesssim 20$ ms. Throughout, an 18° total phase error is used as the threshold for retaining $\geq 90\%$ of ideal coherent gain.

2 Contributions

1. A fully decentralized joint frequency-*and*-phase consensus algorithm (DFPC), extending prior decentralized work that aligned frequency only, with convergence to the arithmetic average of initial states.
2. A closed-form steady-state total phase error for DFPC (Eq. 27), derived by propagating per-iteration error injections through the consensus recursion, with the graph dependence captured through λ_2 .

3. A practical, parameterized error model connecting oscillator data-sheet quantities (Allan deviation coefficients, integrated phase-noise power) and estimator quality (CRLBs vs. L and SNR) to array-level phase coherence, including the effect of TDMA-shared observation time.
4. KF-DFPC: the integration of per-node Kalman filtering with consensus averaging, including the covariance bookkeeping across the mixing step, at per-node complexity $\mathcal{O}(\text{card}\{\chi_n\} + 8)$ —no asymptotic overhead relative to plain DFPC for dense graphs.
5. A thorough comparative simulation study (DFPC, KF-DFPC, DLMS, DKF, KCIF) across array size, connectivity, SNR, and update interval, establishing when Kalman filtering pays: short update intervals, low SNR, and sparse connectivity.

3 Limitations

The following are assessed from a full read of the published text.

1. **No propagation channel in the measurement model.** Nodes observe each other’s oscillator signal directly in AWGN (their Eq. 15). All geometry- and channel-induced phases are lumped into a uniform constant θ_n^e assumed corrected elsewhere. The identifiability problem this hides—a one-way receiver observes oscillator *plus* channel phase and cannot separate them without reciprocity or calibration—is central to any over-the-air realization and is not addressed. Multipath appears only as a qualitative remark.
2. **Statistics-level validation.** The simulations inject Gaussian draws from the assumed error distributions (Algorithm 1) rather than processing waveforms through a receiver. The theory is therefore validated against its own assumptions; estimator failures (missed detection, timing error, cycle slips), quantization, and analog impairments are outside the model.
3. **Instantaneous, idealized actuation.** Consensus updates are applied as immediate state assignments each iteration. There is no correction latency, no NCO quantization, and no control-loop dynamics; the single mention of latency concerns CRLB observation windows. Consequently the analysis cannot expose the error contribution of estimator uncertainty propagated over an actuation delay, which dominates real closed-loop implementations at short intervals.
4. **No closed form for the Kalman variant.** Eq. 27 covers DFPC only; KF-DFPC’s steady state is characterized purely by simulation against the DFPC bound. The steady-state (Riccati) error floor of the filtered loop is not derived.
5. **Oscillator model omissions.** The two-term Allan-deviation law captures white FM and random-walk FM but not flicker FM; phase jitter is drawn independently each interval rather than evolving as a continuous process, so intra- and inter-interval phase noise are statistically disconnected.
6. **Idealized coordination plane.** Time synchronization is assumed complete; sample-clock offset and its accumulated timing drift are absent; the local broadcasts of estimates are error-free and instantaneous.

4 Relevance to this project

This is the closest published prior art to the error-floor analysis pursued in this repository. Its Eq. 27 anchors the framing: two of its terms (drift $\times$ interval and estimation error) reproduce terms of our decomposition, so any contribution must be stated as what Eq. 27 lacks — actuation latency as a parameter distinct from the update interval, coupled to the Kalman steady-state posterior; an analytical flicker term; a continuous phase-noise process; and validation by an impairment-complete sampled-IQ simulation rather than draws from the assumed statistics. The paper’s 18°/90%-gain threshold and its error-vs-update-interval design chart (their Fig. 10) are the natural presentation conventions to adopt and extend. Its consensus layer is complementary rather than competing: it operates above an idealized pairwise exchange, which is precisely the physical link our simulator models.